



CUDA be good

Conor Williams


August 23, 2023

Content:

```
for title in ["Hard LSTM", "Multi-blank"]:
```


- Background
- Optimization and implementation
- Results
- Further work

LSTM background




Source: <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

LSTM background



Source: <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>


LSTM background



Source: <https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

State of affairs:

- Pytorch.
- CuDNN.
- Only *soft* activation functions.
- Pure python 3× slower (no autocast).



The challenge:

- Implement a from-scratch LSTM.
- *Correct* forward and backward.
- As fast as CuDNN's hand optimized assembly.

Optimizations:


Goal: spend all our time in (big) GEMMs.



Optimizations:

Goal: spend all our time in (big) GEMMs.

- Fuse GEMM across time, let $X_{ab} = x_a^b$:




Optimizations:

Goal: spend all our time in (big) GEMMs.

- Fuse GEMM across time, let $X_{ab} = x_a^b$:

$$\{Wx^0, Wx^1, Wx^2, \dots\} \rightarrow W \begin{bmatrix} x^0 & x^1 & \dots \\ \downarrow & \downarrow & \downarrow \end{bmatrix}$$




Optimizations:

Goal: spend all our time in (big) GEMMs.

- Fuse GEMM across time, let $X_{ab} = x_a^b$:

$$\{Wx^0, Wx^1, Wx^2, \dots\} \rightarrow W \begin{bmatrix} x^0 & x^1 & \dots \\ \downarrow & \downarrow & \downarrow \end{bmatrix}$$

- Recurrent steps \rightarrow CUDA.




Optimizations:

Goal: spend all our time in (big) GEMMs.

- Fuse GEMM across time, let $X_{ab} = x_a^b$:

$$\{Wx^0, Wx^1, Wx^2, \dots\} \rightarrow W \begin{bmatrix} x^0 & x^1 & \dots \\ \downarrow & \downarrow & \downarrow \end{bmatrix}$$

- Recurrent steps \rightarrow CUDA.
- Fuse weights and bias.




Optimizations:

Goal: spend all our time in (big) GEMMs.

- Fuse GEMM across time, let $X_{ab} = x_a^b$:

$$\{Wx^0, Wx^1, Wx^2, \dots\} \rightarrow W \begin{bmatrix} x^0 & x^1 & \dots \\ \downarrow & \downarrow & \downarrow \end{bmatrix}$$

- Recurrent steps \rightarrow CUDA.
- Fuse weights and bias.
- Single activation kernel.




CuBLAS:

```
for t in sequence:
    gates[t] = y[t] @ R + x[t] @ W + b # Pre-computing

    state = kernel(gates[t]) # CUDA

    y_out[t] = state.out
    y[t + 1] = state.next
```



LSTM results:




LSTM results:

Hard activation functions are now $7.7\times$ faster than the legacy code!


RNN-T inference:




RNN-T inference:




RNN-T inference:




RNN-T inference:




RNN-T inference:




RNN-T inference:




RNN-T inference:




RNN-T inference:




RNN-T inference:




RNN-T inference:




Multi-blank inference:




Multi-blank inference:




Multi-blank inference:




Multi-blank inference:




Multi-blank training:




Multi-blank training:




Multi-blank training:



Multi-blank training:



Multi-blank training:



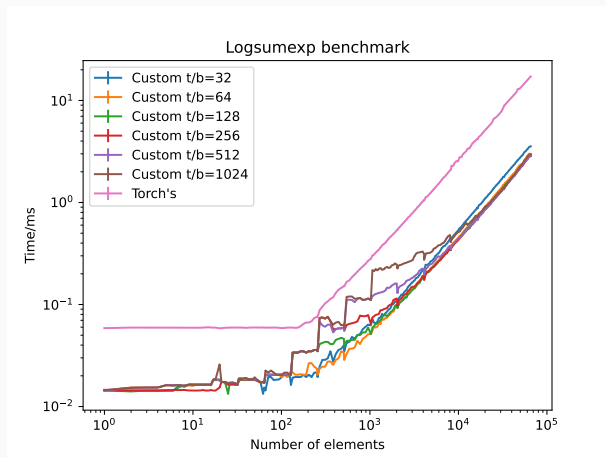
Log softmax:

Let L be the log_softmax function $L: \mathbb{R}^n \rightarrow \mathbb{R}^n$.

$$\begin{aligned} L(x)_i &= \log \frac{\exp x_i}{\sum_j \exp x_j} \\ &= x_i - \log \sum_j \exp x_j \\ &= x_i - \alpha - \log \sum_j \exp (x_j - \alpha) \\ &= x_i - f(x) \end{aligned}$$



Log sum of exponentials:



Multi-blank results:

- Half memory of the legacy code → 85% larger batches.
- 1.2× faster with no big-blanks.
- Supports multi-blank.
- Unoptimised MB WER same as no MB.
- 1.35× GPU inference speedup (unoptimized).

Extensions:

- LSTM train with hard activation functions.
- Tune muliblack hyperparameters, e.g. undernormalization, big-blank set.
- Optimize torches map-reduce operations when in-place.
- FastEmit.